

Occlusion Culling Algorithms

Developers are always going to want better performance in real-time rendering, and so speed-up techniques and acceleration schemes will always be needed. In this excerpt from Chapter 7, "Speed-Up Techniques," of *Real-Time Rendering*, the authors discuss the class of acceleration schemes known as the occlusion culling techniques.

by Tomas Möller



**Excerpted from *Real-Time Rendering*
(AK Peters, 1999)**



One of the great myths concerning computers is that one day we will have enough processing power to do everything. Even in a relatively simple application like word processing, we find that additional power is always needed. This power is applied to all sorts of things, such as on-the-fly spell and grammar checking, more elaborate text formatting, better presentation, antialiased text display, automated voice recognition and dictation, etc.

In real-time rendering we have at least three performance goals: more frames per second, higher resolution, and more (and more realistic) objects in the scene. A speed of 60–72 frames per second is generally considered enough, and perhaps 1600x1200 is enough resolution for a while, but there is no real upper limit on scene complexity. The rendering of a Boeing-777 would include 132,500 parts and over 3,000,000 fasteners, which would yield a polygonal model with over 500,000 polygons [Cripe98]. Our conclusion: speed-up techniques and acceleration schemes will always be needed.

In this article, we will talk about a certain class of acceleration scheme called *occlusion culling techniques*. Most of this article is an excerpt from chapter 7, "Speed-Up Techniques" from *Real-Time Rendering* (www.realtimerendering.com or www.acm.org/tog/resources/RTR/). The occlusion culling section is preceded by sections on backface and clustered culling, hidden surface removal, view-frustum culling, portal culling, and detail culling. Sections on impostor algorithms, level of detail techniques, triangle fan, strip and polygon mesh techniques follow after.

To *cull* can mean to "select from a flock," and in the context of computer graphics this is exactly what *culling techniques* do. The flock is the whole scene that we want to render, and the selection is made to those portions of the scene that are not considered to contribute to the final image. The

smart mechanism in all respects. For example, it has the following implications. Imagine the viewer is looking along a line where 10 spheres are placed. This is illustrated in Figure 1.




Figure 1. An illustration of how occlusion culling can be useful. Ten spheres are placed in a line, and looking along this line (left). The depth complexity image in the middle shows that some pixels are covered several times, even though the final image (on the right) only shows one sphere.

An image rendered from this viewpoint will show but one sphere, even though all 10 spheres have been scan-converted and compared to the Z-buffer and then potentially written to the color buffer. The simple conclusion in this case is that nine spheres will be drawn unnecessarily. This kind of uninteresting scene is not that likely to be found in reality, but it describes (from its viewpoint) a densely populated model. These sorts of configurations are found in such real scenes as a car engine, a city, and the inside of a skyscraper.

Thus it seems plausible that an algorithmic approach to avoid this kind of inefficiency makes sense in terms of speed. Such approaches go under the name of *occlusion culling algorithms*, since they cull away (avoid drawing) objects that are occluded, that is, inside the view frustum but not visible in the final image. The optimal occlusion culling algorithm would select only the objects that are visible. In a sense, the Z-buffer selects and renders only those objects which are visible, but not all objects are sent through the pipeline. The idea behind efficient occlusion culling algorithms is to perform some simple tests early on and so avoid sending data through much of the pipeline.

Pseudocode for a general occlusion culling algorithm is shown in Figure 2, where the function `isOccluded`, often called the *visibility test*, checks whether an object is occluded. *G* is the set of geometrical objects to be rendered; *OR* is the occlusion representation.

```
1: OcclusionCullingAlgorithm (G)
2: OR = empty
3: for each object g in G
4:   if (isOccluded(g, OR))
5:     Skip(g)
6:   else
7:     Render(g)
8:     Update(OR, g)
9: end
10: end
```

updated with that object.

For some algorithms, it is expensive to update the occlusion representation, so this is only (before the actual rendering starts) with the objects that are believed to be good occluders, and then updated from frame to frame.

A number of occlusion algorithms will be scrutinized in this section.

Hierarchical Z-Buffering and the Hierarchical Visibility Algorithm

One approach to occlusion culling is the *hierarchical visibility* (HV) algorithm [Greene93]. This algorithm maintains the scene model in an octree, and a frame's Z-buffer as an image pyramid, which we call a Z-pyramid. The octree enables hierarchical culling of occluded regions of the scene, and the Z pyramid enables hierarchical Z-buffering of individual primitives and bounding volumes. The Z-pyramid is thus the occlusion representation of this algorithm. Examples of these data structures are shown in Figure 3.

Any method can be employed for organizing scene primitives in an octree, although Greene [Greene93] recommend a specific algorithm that avoids assigning small primitives to large nodes. In general, an octree is constructed by enclosing the entire scene in a minimal axis-aligned bounding box. The rest of the procedure is recursive in nature, and starts by checking whether the box contains fewer than a threshold number of primitives. If it does, the algorithm binds the primitives to the node and then terminates the recursion. Otherwise, it subdivides the box along its main axes using planes, thereby forming eight boxes (hence the name octree). Each new box is tested and subdivided again into 2x2x2 smaller boxes. This process continues until each box contains at most the threshold number of primitives, or until the recursion has reached a specified deepest level [Samet89a,Samet89b]. This is illustrated in two dimensions, where the data structure is called a *quadtree*, in Figure 4.




Figure 3. Example of occlusion culling with the hierarchical visibility algorithm [Greene95], showing the scene (lower right) with the corresponding Z-pyramid (on the left), and octree subdivision (upper right). The algorithm traverses the octree from front to back and culling occluded octree nodes as they are encountered. The algorithm only visits visible octree nodes and their children (the nodes portrayed at the upper right). The algorithm renders the polygons in visible boxes. In this example, culling of occluded octree nodes reduces the

to process the contents of that box further, since its contents do not contribute to the final image. Otherwise, we render the primitives associated with the node into the Z-pyramid (*tileInto* pseudocode) and then process each of the node's children (if it has any) in front-to-back using this same recursive procedure. When recursion finishes, all visible primitives have been tiled into the pyramid, and a standard Z-buffer image of the scene has been created.




Figure 4. The construction of a quadtree (which is the two-dimensional version of an octree). The construction starts from the left by enclosing all objects in a bounding box. Then the boxes are recursively divided into equal-sized boxes until each box (in this case) is empty or contains one object.

The HV algorithm performs occlusion culling very efficiently because it only traverses visible nodes and their children, and it only renders the primitives in visible nodes. This can save a lot of work in scenes that are densely occluded. For example, in the scene pictured in Figure 3, more than 99% of on-screen polygons are inside occluded octree nodes, which are therefore culled before being rendered to the pyramid [Greene95].

```

1: ProcessOctreeNode(OctreeNode N)
2: if(isOccluded(NBV, ZP)) then return;
3: for each primitive p in N
4:   tileInto(p, ZP)

```

```

7: ProcessOctree(C)
8: end

```

Figure 5. Pseudocode for the hierarchical visibility algorithm. To render a frame this procedure is called on the root node of the octree. NBV is the bounding volume of the octree node *N*, and ZP is the Z-pyramid that is the occlusion representation of this algorithm. The operation *tileInto* renders a primitive *p* into the Z-pyramid and this also updates the entire Z-pyramid.

Now we will describe how the Z-pyramid is maintained and how it is used to accelerate occlusion culling. The finest (highest-resolution) level of the Z-pyramid is simply a standard Z-buffer. At all other levels, the z-value is the farthest z in the corresponding 2x2 window of the adjacent finer level. There, the z-value represents the farthest z for a square region of the screen. To maintain a Z-pyramid, when a z-value is overwritten in the Z-buffer it is propagated through the coarser levels of the Z-pyramid.

Figure 6. On the left, a 4x4 piece of the Z-buffer is shown. The numerical values are the actual z-values downsampled to a 2x2 region where each value is the farthest (largest) of the four 2x2 regions on the right. The farthest value of the remaining four z-values is computed. These three maps compose an image which is called the hierarchical Z-buffer.

Next, we describe how hierarchical culling of octree nodes is done. To determine whether a node is visible, the front faces of its bounding box are tested against the Z-pyramid. The node is considered visible if any of its front faces are not occluded by the Z-pyramid. To establish whether an individual face is visible, we begin at the coarsest Z-pyramid cell that encloses the face's screen projection. The face's depth within the cell

(*znear*) is then compared to the Z-pyramid value, and if *znear* is farther, the face is known to be occluded. For densely occluded scenes, this procedure often culls an occluded face with a single depth comparison. When this initial test fails to cull a face, its visibility can be definitively established by recursively traversing from the initial Z-pyramid cell down to finer levels in the Z-pyramid. Additional depth comparisons at the encountered child cells must be performed. If this search procedure does not ultimately find a visible sample on the face, the face is occluded. In scenes where the bounding boxes of octree nodes overlap deeply on the screen, this hierarchical procedure can establish visibility much more efficiently than can conventional Z-buffering. For example, the scene pictured in Figure 3, hierarchical culling of octree nodes with the Z-pyramid generates roughly one hundred times fewer depth comparisons than visibility testing with conventional Z-buffering.

As we have seen, the original HV algorithm [Greene93] used a Z-pyramid only for occlusion testing and employed traditional Z-buffer scan conversion to render the polygons in visible octree nodes. Subsequently, a more efficient method, called *hierarchical polygon tiling* [Greene96], was developed. This algorithm adapts the basic screen subdivision procedure described above to polygons. The goal is, finding the visible samples on a polygon. When tiling a polygon into the Z-pyramid, it must be determined, at each level of subdivision where the polygon is compared to a Z-pyramid cell, whether the polygon overlaps the cell and if so, whether the polygon is occluded by the cell. Overlap testing is accelerated with coverage masks that indicate which subcells within a Z-pyramid cell are covered by the polygon, and occlusion tests are performed by comparing the polygon's nearest z-value to the Z-pyramid value. This procedure for hierarchical Z-buffering is very efficient, because it only traverses regions of the screen where a polygon is visible or nearly visible.

Hierarchical polygon tiling can be performed even more efficiently if the polygons in a scene are organized into a BSP tree or an "octree of BSP trees" [Greene96]. The reason is that this enables traversing polygons in strict front-to-back order, which eliminates the need to maintain occlusion information. Rather, the only occlusion information required is whether or not each image pixel has been written, and this information can be maintained in an image pyramid of coverage masks, called a *coverage pyramid*. This is the data structure maintained by hierarchical polygon tiling with coverage masks [Greene96], which is similar to hierarchical Z-buffering except that occlusion tests are performed with coverage-mask operations instead of depth comparisons, which accelerates the process considerably.

The HV algorithm implemented with hierarchical tiling may well be the most efficient method for software rendering of complex scenes composed of polygons, but it is not fast enough for real-time rendering of complex scenes on today's microprocessors. To enable real-time rendering, et al. [Greene93] suggest modifying hardware Z-buffer pipelines to support HV, which requires

pass, the standard HV algorithm is run in software, traversing the octree from front to back, skipping nodes which have already been rendered. This second pass fills in any missing polygons in the scene. The final step in processing a frame is to update the visible node list. Typically, this HV algorithm runs considerably faster than the all-software version, because nearly all polygons are rendered with Z-buffer hardware.

Greene and Kass [Greene94b] have developed an extension to hierarchical Z-buffering for antialiased scenes with error bounds. Another interesting algorithm for occlusion culling is the skeleton developed by Durand et al. [Durand97,Durand97b].

The HOM Algorithm

The *hierarchical occlusion map* (HOM) algorithm [Zhang97] is another way of enabling hierarchical image-space culling (such as the hierarchical Z-buffering algorithm). However, the HOM can be used on systems that have graphics hardware but not a hardware Z-pyramid, and can handle dynamic scenes. The HOM algorithm is described in detail in Zhang's Ph.D. thesis [Zhang97].

We start by describing how the function `isOccluded` works. This function, used in the pseudocode in Figure 2, is a key part of the algorithm. This occlusion test takes place after the view transform and the viewer is located at the origin looking down the negative z-axis, with the x-axis going to the right and the y-axis going upwards. The test is then divided into two parts: a one-dimensional *depth test* in the z-direction and a two-dimensional *overlap test* in the xy plane, i.e., whereby the image generated by the occluders is compared to the image of the object being tested. The overlap test supports approximate visibility culling, where objects that "shine through" the occluders can be culled away using an opacity threshold parameter.

For both tests, a set of potentially good occluders is identified before the scene is rendered. An occlusion representation is built from these. This step is followed by the rendering of the scene. First the occluders are rendered without an occlusion test. Then the rest of the scene is processed, having each object tested against the occlusion representation. If the object is occluded by the representation, it is not rendered.

For the two-dimensional overlap test, the occluders are first rendered into the color buffer. The color is stored on a black background. Therefore, texturing, lighting, and Z-buffering can be turned off. The advantage of this operation is that a number of small occluders can be combined into a single occluder. The

rendered image, which is called an *occlusion map*, is read back into the main memory of the computer. For simplicity, we assume that this image has the resolution of $2^n \times 2^n$ pixels, where n is the base for the occlusion representation. Then a *hierarchy of occlusion maps* (HOM), i.e., a pyramid of occlusion maps, is created by averaging over $2^{n-1} \times 2^{n-1}$ pixel blocks to form a map of $2^{n-1} \times 2^{n-1}$ pixels. This is done recursively until a minimum size is reached (for example, 16×16 pixels). The highest-resolution level of the HOM is numbered 0, with increasing numbers having decreasing resolution. The gray-scale values in the HOM are said to be the *opacity* of the pixels. A high value (near white) for a pixel at a level above 0 means that most of the pixels it represents are covered by the HOM.

The creation of the HOM can be implemented either on the CPU or by texture mapping, with bilinear interpolation used as a minification filter. For large image sizes, the texture filtering approach was found to be faster, and for small image sizes, the CPU was faster. Of course, this varies with the graphics hardware. For a 1024×1024 image, Zhang et al. [Zhang97] used a 256×256 image for the HOM. An example of a HOM is shown in Figure 7.

The overlap test against the HOM starts by projecting the bounding volume of the object onto the screen (Zhang et al. [Zhang97] used oriented bounding boxes). This projection is bounded by a rectangle, which then covers more pixels than the object enclosed in the bounding volume. So this test is a *conservative test*, meaning that even if the test results show that an object is not occluded, it may still be so. This rectangle is then compared against the HOM for overlap. The overlap test starts at the level in which the size of the pixel in the HOM is approximately the size of the rectangle. If all pixels in the rectangle are opaque (which means fully white for non-approximate visibility culling), then the rectangle is occluded in the xy plane and the object is said to pass the test. On the other hand, if a pixel is not opaque, then the test for that pixel continues recursively to the next level in the HOM which are covered by the rectangle, meaning that the resolution of the occlusion map increases with each test.

For approximate visibility culling, the pixels in the HOM are not compared to full opacity, i.e. white, but rather against an opacity threshold value, a gray-scale value. The lower the threshold value, the more approximate the culling. The advantage here is that if a pixel is not fully opaque (white) but above the threshold, then the overlap test can terminate earlier. The penalty is that some objects are omitted from rendering even though it is (partially) visible. The opacity values are not compared from one level to another in the HOM, as shown in the following example.

Example: *Computation of opacity threshold values*

Assume the rendered image is 1024x1024 pixels and that the lowest level in the HOM (i.e., the largest resolution) has a resolution of 128x128 pixels. A pixel in this level-zero occlusion map corresponds to an 8x8-pixel region in the rendered image. Also assume that a 2x2 region of pixels in an 8x8 region can pass as a negligible hole. This would give an opacity value $O_0 = 1 - 2^2/8^2 = 0.9375$. The next level in the HOM would then have a 64x64 resolution, and a pixel in this level would correspond to 16x16 pixels in the rendered image. So the opacity threshold at this level is $O_1 = 1 - 2^2/16^2$ which is approximately 0.984.

We will now derive a recursive formula for computing the opacity values of the different levels in the HOM. The opacity of the level with the highest resolution in the HOM is $O_0 = 1 - n/m$, where n is the number of black pixels that can be considered a negligible hole, and m is the number of pixels in the rendered image represented by one pixel in this occlusion map ($m = 8 \times 8$ in the example). The next level in the HOM has a threshold of $O_1 = 1 - n/(4m) = 1 - (1 - O_0)/4 = (3 + O_0)/4$. This reasoning can be generalized to the formula in below for the k th level in the HOM.

$$O_{k+1} = (3 + O_k)/4$$

For more details on this topic, consult Zhang's Ph.D. thesis [Zhang98].

For the one-dimensional z-depth test, we must be able to determine whether an object is behind any of the selected occluders. Zhang [Zhang98] describes a number of methods, and we choose to use the *depth estimation buffer*, which provides reasonable estimation and does not require a Z-buffer. It is implemented as a software Z-buffer that divides the screen into a number of rectangular regions. These regions are rather large in relation to the pixel size. The selected occluders are inserted into this buffer. In each region the farthest z-value is stored. This is in contrast to a normal Z-buffer, which stores the z-value for each pixel.

Figure 8. Illustration of the depth estimation buffer. Illustration after Zhang [Zhang98].

The depth estimation buffer is built for each frame. During rendering, to test whether an object is behind the occluders, the depth test (i.e., whether it is behind the occluders) the z-value of the nearest vertex of the object's bounding box is computed. This value is compared against the z-values of all regions in the depth estimation buffer that the bounding box rectangle covers in screen space. If the near value of the bounding box is larger than the stored z-depth in all regions, then the object passes the depth test, and is not occluded in the depth direction. A resolution of 64x64 regions in the depth estimation buffer is used by Zhang et al. [Zhang97].

For an object to be occluded, it must thus first pass the overlap test; i.e., the rectangle of the object's bounding volume of the object must pass the HOM test. Then it must pass the depth test, i.e., be behind the occluders. If an object passes both tests, the object is occluded and is not rendered.

Before the algorithm starts, an occluder database is built, where a few criteria are used to select certain objects [Zhang98]. First, small objects do not work well in the occluder database, as they usually cover a small portion of the image unless the viewer is very close to them. Second, objects with a high polygon count are also avoided, as these may negatively affect the performance of the rendering of the occlusion map. Third, objects with large or ill-shaped bounding boxes (e.g., a long bounding box for a skinny polygon) should be avoided, as they may cause the depth estimation to be too conservative. Finally, redundant objects are avoided: for example, a clock on a wall does not contribute as much to occlusion as the wall itself.

At runtime, occluders are selected from the database. To avoid allowing the creation of too many occluders to become a bottleneck, there is a limit to the number of occluders that can be selected. The selection is done with respect to their distance from the viewer and to their size. Only objects inside the viewer's frustum are selected. A case when this does not work well is shown in Figure 9.

Get daily news, dev blogs, and stories from Game Developer straight to your inbox

[Subscribe](#)

[LATEST JOBS](#)

[TECHNICAL LIGHTING
ARTIST – TURN 10
STUDIOS](#)

[SR CHARACTER ARTIST –
GEARS OF WAR – THE
COALITION](#)

[SOCIAL MEDIA /
COMMUNITY MANAGER](#)

[MORE JOBS](#) 

CONNECT WITH US

Explore the
Game Developer Job Board

Browse open positions across the game
industry or recruit new talent for your
studio

[BROWSE](#)

Subscribe to
Game Developer Newsletter

Get daily Game Developer top stories
every morning straight into your inbox

[SUBSCRIBE](#)

Foll
@g

Foll
to-c
info

Discover More From Informa Tech

[Game Developers Conference](#)

[Game Developer Jobs](#)

[GDC Vault](#)

[Game Career Guide](#)

[Independent Games Festival](#)

[Game Developers Choice Awards](#)

[Omdia](#)

Working With Us

[Contact Us](#)

[About Us](#)

[Advertise](#)

Fo



Get daily news, dev blogs, and stories from Game Developer straight to your inbox

Subscribe

LATEST JOBS

Microsoft

REDMOND, WA, USA

1.04.23

Microsoft

HYBRID (VANCOUVER, BC,
CANADA)

1.04.23

Fast Travel Games

HYBRID (STOCKHOLM, SWEDEN)

1.09.23

Nig

LOS A

1.09.

**TECHNICAL LIGHTING
ARTIST – TURN 10
STUDIOS**

**SR CHARACTER ARTIST
– GEARS OF WAR – THE
COALITION**

**SOCIAL MEDIA /
COMMUNITY MANAGER**

**LEV
SCR
STU**

MORE JOBS 

CONNECT WITH US

Explore the

Subscribe to

Follow us

Discover More From Informa Tech

[Game Developers Conference](#)

[Game Developer Jobs](#)

[GDC Vault](#)

[Game Career Guide](#)

[Independent Games Festival](#)

[Game Developers Choice Awards](#)

[Omdia](#)

Working With Us

[Contact Us](#)

[About Us](#)

[Advertise](#)

Follow Game Developer



[Home](#)

[Cool](#)

Copyright © 2023 Informa PLC Informa UK Limited is a company registered in England and Wales
1072954 whose registered office is 5 Howick

